

# Assignment\_13(Try\_various\_CNN\_networks\_on\_MNIST\_dataset)

September 18, 2018

## 1 OBJECTIVE :- Try various CNN networks on MNIST dataset

```
In [1]: # Importing libraries
        from __future__ import print_function
        import keras
        from keras.datasets import mnist
        from keras.models import Sequential
        from keras.layers import Dense, Dropout, Flatten
        from keras.layers import Conv2D, MaxPooling2D
        from keras import backend as K
        from keras.initializers import he_normal
        from keras.layers.normalization import BatchNormalization
        import numpy as np
        import matplotlib.pyplot as plt
        %matplotlib inline

        batch_size = 128
        num_classes = 10
        epochs = 12

        # input image dimensions
        img_rows, img_cols = 28, 28

        # the data, split between train and test sets
        (x_train, y_train), (x_test, y_test) = mnist.load_data()
```

Using TensorFlow backend.

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>  
11493376/11490434 [=====] - 2s 0us/step

```
In [2]: if K.image_data_format() == 'channels_first':
        x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
        x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
```

```

        input_shape = (1, img_rows, img_cols)
    else:
        x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
        x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
        input_shape = (img_rows, img_cols, 1)

    x_train = x_train.astype('float32')
    x_test = x_test.astype('float32')
    x_train /= 255
    x_test /= 255
    print('x_train shape:', x_train.shape)
    print(x_train.shape[0], 'train samples')
    print(x_test.shape[0], 'test samples')

    # convert class vectors to binary class matrices
    y_train = keras.utils.to_categorical(y_train, num_classes)
    y_test = keras.utils.to_categorical(y_test, num_classes)

```

```

x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples

```

```

In [0]: # this function is used draw Categorical Crossentropy Loss VS No. of epochs plot
def plt_dynamic(x, vy, ty):
    plt.figure(figsize=(10,5))
    plt.plot(x, vy, 'b', label="Validation Loss")
    plt.plot(x, ty, 'r', label="Train Loss")
    plt.xlabel('Epochs')
    plt.ylabel('Categorical Crossentropy Loss')
    plt.title('\nCategorical Crossentropy Loss VS Epochs')
    plt.legend()
    plt.grid()
    plt.show()

```

## 1.1 (1). CNN with 3 Convolutional layers and kernel size - (3X3)

```

In [5]: # Initialising the model
model_3 = Sequential()

# Adding first conv layer
model_3.add(Conv2D(32, kernel_size=(3, 3),activation='relu',input_shape=input_shape))

# Adding second conv layer
model_3.add(Conv2D(64, (3, 3), activation='relu'))

# Adding Maxpooling layer
model_3.add(MaxPooling2D(pool_size=(2, 2)))

```

```

# Adding Dropout
model_3.add(Dropout(0.25))

# Adding third conv layer
model_3.add(Conv2D(128, (3, 3), activation='relu'))

# Adding Maxpooling layer
model_3.add(MaxPooling2D(pool_size=(2, 2)))

# Adding Dropout
model_3.add(Dropout(0.25))

# Adding flatten layer
model_3.add(Flatten())

# Adding first hidden layer
model_3.add(Dense(256, activation='relu', kernel_initializer=he_normal(seed=None)))

# Adding Dropout
model_3.add(Dropout(0.5))

# Adding output layer
model_3.add(Dense(num_classes, activation='softmax'))

# Printing model Summary
print(model_3.summary())

# Compiling the model
model_3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Fitting the data to the model
history_3 = model_3.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1

```

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 26, 26, 32)	320
conv2d_5 (Conv2D)	(None, 24, 24, 64)	18496
max_pooling2d_3 (MaxPooling2D)	(None, 12, 12, 64)	0
dropout_4 (Dropout)	(None, 12, 12, 64)	0
conv2d_6 (Conv2D)	(None, 10, 10, 128)	73856
max_pooling2d_4 (MaxPooling2D)	(None, 5, 5, 128)	0

dropout_5 (Dropout)	(None, 5, 5, 128)	0
flatten_2 (Flatten)	(None, 3200)	0
dense_3 (Dense)	(None, 256)	819456
dropout_6 (Dropout)	(None, 256)	0
dense_4 (Dense)	(None, 10)	2570

Total params: 914,698  
 Trainable params: 914,698  
 Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 222s 4ms/step - loss: 0.2171 - acc: 0.9305 - va

Epoch 2/12

60000/60000 [=====] - 240s 4ms/step - loss: 0.0674 - acc: 0.9788 - va

Epoch 3/12

60000/60000 [=====] - 237s 4ms/step - loss: 0.0496 - acc: 0.9850 - va

Epoch 4/12

60000/60000 [=====] - 237s 4ms/step - loss: 0.0412 - acc: 0.9880 - va

Epoch 5/12

60000/60000 [=====] - 219s 4ms/step - loss: 0.0357 - acc: 0.9896 - va

Epoch 6/12

60000/60000 [=====] - 209s 3ms/step - loss: 0.0319 - acc: 0.9901 - va

Epoch 7/12

60000/60000 [=====] - 218s 4ms/step - loss: 0.0274 - acc: 0.9912 - va

Epoch 8/12

60000/60000 [=====] - 209s 3ms/step - loss: 0.0243 - acc: 0.9923 - va

Epoch 9/12

60000/60000 [=====] - 205s 3ms/step - loss: 0.0227 - acc: 0.9927 - va

Epoch 10/12

60000/60000 [=====] - 213s 4ms/step - loss: 0.0210 - acc: 0.9932 - va

Epoch 11/12

60000/60000 [=====] - 206s 3ms/step - loss: 0.0201 - acc: 0.9935 - va

Epoch 12/12

60000/60000 [=====] - 208s 3ms/step - loss: 0.0190 - acc: 0.9941 - va

In [6]: # Evaluating the model

score = model\_3.evaluate(x\_test, y\_test, verbose=0)

print('Test score:', score[0])

print('Test accuracy:', score[1])

```

# Test and train accuracy of the model
model_3_test = score[1]
model_3_train = max(history_3.history['acc'])

# Plotting Train and Test Loss VS no. of epochs
# list of epoch numbers
x = list(range(1,epochs+1))

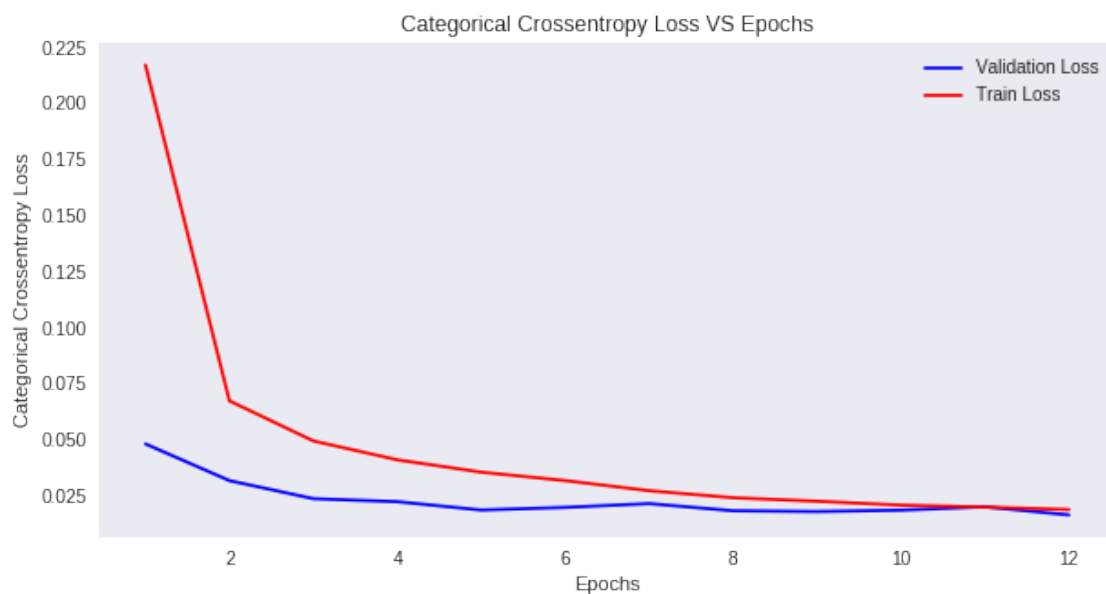
# Validation loss
vy = history_3.history['val_loss']
# Training loss
ty = history_3.history['loss']

# Calling the function to draw the plot
plt_dynamic(x, vy, ty)

```

Test score: 0.01658615784635249

Test accuracy: 0.994



## 1.2 (2). CNN with 5 Convolutional layers and kernel size - (5X5)

In [10]: *# Initialising the model*

```
model_5 = Sequential()
```

*# Adding first conv layer*

```
model_5.add(Conv2D(8, kernel_size=(5, 5),padding='same',activation='relu',input_shape=
```

```

# Adding second conv layer
model_5.add(Conv2D(16, (5, 5), activation='relu'))

# Adding Maxpooling layer
model_5.add(MaxPooling2D(pool_size=(2, 2),padding='same'))

# Adding Dropout
model_5.add(Dropout(0.25))

# Adding third conv layer
model_5.add(Conv2D(32, (5, 5),padding='same', activation='relu'))

# Adding Maxpooling layer
model_5.add(MaxPooling2D(pool_size=(2, 2),padding='same'))

# Adding Dropout
model_5.add(Dropout(0.25))

# Adding fourth conv layer
model_5.add(Conv2D(64, (5, 5),padding='same',activation='relu'))

# Adding fifth conv layer
model_5.add(Conv2D(64, (5, 5), activation='relu'))

# Adding Maxpooling layer
model_5.add(MaxPooling2D(pool_size=(2, 2),padding='same'))

# Adding Dropout
model_5.add(Dropout(0.25))

# Adding flatten layer
model_5.add(Flatten())

# Adding first hidden layer
model_5.add(Dense(256, activation='relu',kernel_initializer=he_normal(seed=None)))

# Adding Batch Normalization
model_5.add(BatchNormalization())

# Adding Dropout
model_5.add(Dropout(0.5))

# Adding output layer
model_5.add(Dense(num_classes, activation='softmax'))

# Printing model Summary
print(model_5.summary())

```

```

# Compiling the model
model_5.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Fitting the data to the model
history_5 = model_5.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1)

```

Layer (type)	Output Shape	Param #
conv2d_20 (Conv2D)	(None, 28, 28, 8)	208
conv2d_21 (Conv2D)	(None, 24, 24, 16)	3216
max_pooling2d_11 (MaxPooling)	(None, 12, 12, 16)	0
dropout_13 (Dropout)	(None, 12, 12, 16)	0
conv2d_22 (Conv2D)	(None, 12, 12, 32)	12832
max_pooling2d_12 (MaxPooling)	(None, 6, 6, 32)	0
dropout_14 (Dropout)	(None, 6, 6, 32)	0
conv2d_23 (Conv2D)	(None, 6, 6, 64)	51264
conv2d_24 (Conv2D)	(None, 2, 2, 64)	102464
max_pooling2d_13 (MaxPooling)	(None, 1, 1, 64)	0
dropout_15 (Dropout)	(None, 1, 1, 64)	0
flatten_3 (Flatten)	(None, 64)	0
dense_5 (Dense)	(None, 256)	16640
batch_normalization_1 (Batch Normalization)	(None, 256)	1024
dropout_16 (Dropout)	(None, 256)	0
dense_6 (Dense)	(None, 10)	2570

```

Total params: 190,218
Trainable params: 189,706
Non-trainable params: 512

```

```

None
Train on 60000 samples, validate on 10000 samples

```

```

Epoch 1/12
60000/60000 [=====] - 143s 2ms/step - loss: 0.3913 - acc: 0.8772 - va
Epoch 2/12
60000/60000 [=====] - 144s 2ms/step - loss: 0.1050 - acc: 0.9703 - va
Epoch 3/12
60000/60000 [=====] - 139s 2ms/step - loss: 0.0771 - acc: 0.9785 - va
Epoch 4/12
60000/60000 [=====] - 135s 2ms/step - loss: 0.0656 - acc: 0.9815 - va
Epoch 5/12
60000/60000 [=====] - 127s 2ms/step - loss: 0.0560 - acc: 0.9843 - va
Epoch 6/12
60000/60000 [=====] - 136s 2ms/step - loss: 0.0495 - acc: 0.9868 - va
Epoch 7/12
60000/60000 [=====] - 143s 2ms/step - loss: 0.0448 - acc: 0.9879 - va
Epoch 8/12
60000/60000 [=====] - 132s 2ms/step - loss: 0.0424 - acc: 0.9882 - va
Epoch 9/12
60000/60000 [=====] - 129s 2ms/step - loss: 0.0379 - acc: 0.9891 - va
Epoch 10/12
60000/60000 [=====] - 126s 2ms/step - loss: 0.0344 - acc: 0.9902 - va
Epoch 11/12
60000/60000 [=====] - 135s 2ms/step - loss: 0.0310 - acc: 0.9913 - va
Epoch 12/12
60000/60000 [=====] - 138s 2ms/step - loss: 0.0306 - acc: 0.9912 - va

```

```

In [11]: # Evaluating the model
         score = model_5.evaluate(x_test, y_test, verbose=0)
         print('Test score:', score[0])
         print('Test accuracy:', score[1])

         # Test and train accuracy of the model
         model_5_test = score[1]
         model_5_train = max(history_5.history['acc'])

         # Plotting Train and Test Loss VS no. of epochs
         # list of epoch numbers
         x = list(range(1, epochs+1))

         # Validation loss
         vy = history_5.history['val_loss']
         # Training loss
         ty = history_5.history['loss']


         # Calling the function to draw the plot
         plt_dynamic(x, vy, ty)

```

Test score: 0.020986396820967274



Test accuracy: 0.9931



### 1.3 (3). CNN with 7 Convolutional layers and kernel size - (2X2)

```
In [18]: # Initialising the model
model_7 = Sequential()

# Adding first conv layer
model_7.add(Conv2D(32, kernel_size=(2, 2),padding='same',activation='relu',input_shape=(28, 28, 1)))

# Adding second conv layer
model_7.add(Conv2D(32, (2, 2), activation='relu'))

# Adding Maxpooling layer
model_7.add(MaxPooling2D(pool_size=(3, 3), strides=(1,1)))

# Adding Dropout
model_7.add(Dropout(0.3))

# Adding third conv layer
model_7.add(Conv2D(64, (2, 2), activation='relu'))

# Adding Maxpooling layer
model_7.add(MaxPooling2D(pool_size=(2, 2),padding='same'))
```

```

# Adding fourth conv layer
model_7.add(Conv2D(64, (2, 2),padding='same',activation='relu'))

# Adding fifth conv layer
model_7.add(Conv2D(128, (2, 2), activation='relu'))

# Adding Maxpooling layer
model_7.add(MaxPooling2D(pool_size=(3, 3),padding='same'))

# Adding Dropout
model_7.add(Dropout(0.25))

# Adding sixth conv layer
model_7.add(Conv2D(128, (2, 2),padding='same',activation='relu'))

# Adding seventh conv layer
model_7.add(Conv2D(256, (2, 2), activation='relu'))

# Adding Maxpooling layer
model_7.add(MaxPooling2D(pool_size=(2, 2), strides=(1,1)))

# Adding Dropout
model_7.add(Dropout(0.25))

# Adding flatten layer
model_7.add(Flatten())

# Adding first hidden layer
model_7.add(Dense(256, activation='relu',kernel_initializer=he_normal(seed=None)))

# Adding Batch Normalization
model_7.add(BatchNormalization())

# Adding Dropout
model_7.add(Dropout(0.5))

# Adding second hidden layer
model_7.add(Dense(128, activation='relu',kernel_initializer=he_normal(seed=None)))

# Adding Dropout
model_7.add(Dropout(0.25))

# Adding output layer
model_7.add(Dense(num_classes, activation='softmax'))

# Printing model Summary
print(model_7.summary())

```

```

# Compiling the model
model_7.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Fitting the data to the model
history_7 = model_7.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=

```

Layer (type)	Output Shape	Param #
conv2d_60 (Conv2D)	(None, 28, 28, 32)	160
conv2d_61 (Conv2D)	(None, 27, 27, 32)	4128
max_pooling2d_33 (MaxPooling)	(None, 25, 25, 32)	0
dropout_42 (Dropout)	(None, 25, 25, 32)	0
conv2d_62 (Conv2D)	(None, 24, 24, 64)	8256
max_pooling2d_34 (MaxPooling)	(None, 12, 12, 64)	0
conv2d_63 (Conv2D)	(None, 12, 12, 64)	16448
conv2d_64 (Conv2D)	(None, 11, 11, 128)	32896
max_pooling2d_35 (MaxPooling)	(None, 4, 4, 128)	0
dropout_43 (Dropout)	(None, 4, 4, 128)	0
conv2d_65 (Conv2D)	(None, 4, 4, 128)	65664
conv2d_66 (Conv2D)	(None, 3, 3, 256)	131328
max_pooling2d_36 (MaxPooling)	(None, 2, 2, 256)	0
dropout_44 (Dropout)	(None, 2, 2, 256)	0
flatten_8 (Flatten)	(None, 1024)	0
dense_19 (Dense)	(None, 256)	262400
batch_normalization_10 (Batch Normalization)	(None, 256)	1024
dropout_45 (Dropout)	(None, 256)	0
dense_20 (Dense)	(None, 128)	32896
dropout_46 (Dropout)	(None, 128)	0

```

-----
dense_21 (Dense)                (None, 10)                1290
=====

```

```

Total params: 556,490
Trainable params: 555,978
Non-trainable params: 512

```

```

-----
None

```

```

Train on 60000 samples, validate on 10000 samples

```

```

Epoch 1/12
60000/60000 [=====] - 263s 4ms/step - loss: 0.4402 - acc: 0.8525 - va
Epoch 2/12
60000/60000 [=====] - 256s 4ms/step - loss: 0.0911 - acc: 0.9727 - va
Epoch 3/12
60000/60000 [=====] - 262s 4ms/step - loss: 0.0675 - acc: 0.9802 - va
Epoch 4/12
60000/60000 [=====] - 279s 5ms/step - loss: 0.0536 - acc: 0.9839 - va
Epoch 5/12
60000/60000 [=====] - 279s 5ms/step - loss: 0.0446 - acc: 0.9867 - va
Epoch 6/12
60000/60000 [=====] - 282s 5ms/step - loss: 0.0485 - acc: 0.9858 - va
Epoch 7/12
60000/60000 [=====] - 287s 5ms/step - loss: 0.0397 - acc: 0.9882 - va
Epoch 8/12
60000/60000 [=====] - 282s 5ms/step - loss: 0.0382 - acc: 0.9888 - va
Epoch 9/12
60000/60000 [=====] - 281s 5ms/step - loss: 0.0321 - acc: 0.9902 - va
Epoch 10/12
60000/60000 [=====] - 285s 5ms/step - loss: 0.0299 - acc: 0.9913 - va
Epoch 11/12
60000/60000 [=====] - 275s 5ms/step - loss: 0.0288 - acc: 0.9919 - va
Epoch 12/12
60000/60000 [=====] - 281s 5ms/step - loss: 0.0300 - acc: 0.9914 - va

```

```

In [19]: # Evaluating the model
         score = model_7.evaluate(x_test, y_test, verbose=0)
         print('Test score:', score[0])
         print('Test accuracy:', score[1])

         # Test and train accuracy of the model
         model_7_test = score[1]
         model_7_train = max(history_7.history['acc'])

         # Plotting Train and Test Loss VS no. of epochs
         # list of epoch numbers
         x = list(range(1, epochs+1))

```

```


# Validation loss
vy = history_7.history['val_loss']
# Training loss
ty = history_7.history['loss']

# Calling the function to draw the plot
plt_dynamic(x, vy, ty)

```

Test score: 0.026096460982217105

Test accuracy: 0.9921



## 1.4 CONCLUSION

### 1.5 (a). Procedure Followed :

1. Load MNIST dataset
2. Split the dataset into train and test
3. Normalize the train and test data
4. Convert class variable into categorical data vector
5. Implement Softmax classifier with 3, 5 and 7 conv layers .
6. Use kernel -size (3X3) , (5X5) and (2,2) .
7. Draw Categorical Crossentropy Loss VS No.of Epochs plot .

### 1.6 (b) Table (Different models with their train and test accuracies):

```

In [22]: # Installing the library prettytable
!pip install prettytable

```

```

# Creating table using PrettyTable library
from prettytable import PrettyTable

# Names of models
names = ['CNN(3-Conv layers) With Kernel-size = (3,3)', 'CNN(5-Conv layers) With Kernel-size = (5,5)',
         'CNN(7-Conv layers) With Kernel-size = (2,2)']

# Training accuracies
train_acc = [model_3_train, model_5_train, model_7_train]

# Test accuracies
test_acc = [model_3_test, model_5_test, model_7_test]

numbering = [1, 2, 3]

# Initializing prettytable
ptable = PrettyTable()

# Adding columns
ptable.add_column("S.NO.", numbering)
ptable.add_column("MODEL", names)
ptable.add_column("Training Accuracy", train_acc)
ptable.add_column("Test Accuracy", test_acc)

# Printing the Table
print(ptable)

```

Requirement already satisfied: prettytable in /usr/local/lib/python3.6/dist-packages (0.7.2)

S.NO.	MODEL	Training Accuracy	Test Accuracy
1	CNN(3-Conv layers) With Kernel-size = (3,3)	0.99411666666348775	0.994
2	CNN(5-Conv layers) With Kernel-size = (5,5)	0.9913	0.9931
3	CNN(7-Conv layers) With Kernel-size = (2,2)	0.99195	0.9921